

Functions, Attributes, and Class

Data Analysis with R and Python

Deepayan Sarkar



Functions

- Most useful things in R happen by calling functions
- Functions have one or more arguments
 - All arguments have names (available as variables inside the function)
 - Arguments may be compulsory or optional
 - Optional arguments usually have "default" values

Functions

- Most useful things in R happen by calling functions
- Functions have one or more arguments
 - All arguments have names (available as variables inside the function)
 - Arguments may be compulsory or optional
 - Optional arguments usually have "default" values
- Arguments may or may not be named when the function is called
 - Unnamed arguments are matched by position
 - Optional arguments are usually named
- Functions normally also have a useful "return" value

Example: Delhi Air Quality data

```
aqi <- read.csv("https://deepayan.github.io/BSDS/2024-01-DE/data/rkpuram-aqi.csv")  
str(aqi)
```

```
'data.frame':   3930 obs. of  7 variables:  
 $ date: chr  "2024/11/1" "2024/11/2" "2024/11/3" "2024/11/4" ...  
 $ pm25: int  300 306 308 300 282 267 307 275 269 260 ...  
 $ pm10: int  260 249 298 246 227 251 205 198 195 264 ...  
 $ o3 : int  40 57 53 56 52 47 40 46 53 46 ...  
 $ no2 : int  19 16 21 12 15 16 13 9 8 9 ...  
 $ so2 : int  8 5 3 9 8 5 9 10 11 18 ...  
 $ co : int  17 19 17 24 22 19 21 22 19 32 ...
```

Example: Delhi Air Quality data


- How is AQI (PM10) related to `no2` (Nitrogen dioxide)?

Example: Delhi Air Quality data

- How is AQI (PM10) related to `no2` (Nitrogen dioxide)?
- Basic tools we are familiar with
 - Scatter plot
 - Linear Regression


Example: Delhi Air Quality data

```
plot(pm10 ~ no2, data = aqi)
```



Example: Delhi Air Quality data

```
plot(sqrt(pm10) ~ sqrt(no2), data = aqi)
```



Example: Delhi Air Quality data

- Formula for linear regression line

$$\hat{b} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$
$$\hat{a} = \bar{y} - \hat{b}\bar{x}$$

Example: Delhi Air Quality data

- Calculation of linear regression line

```
x <- sqrt(aqi$no2)
y <- sqrt(aqi$pm10)
xbar <- mean(x)
ybar <- mean(y)
sxy <- sum( (x - xbar) * (y - ybar) )
sxx <- sum( (x - xbar)^2 )
```

Example: Delhi Air Quality data

- Calculation of linear regression line

```
x <- sqrt(aqi$no2)
y <- sqrt(aqi$pm10)
xbar <- mean(x)
ybar <- mean(y)
sxy <- sum( (x - xbar) * (y - ybar) )
sxx <- sum( (x - xbar)^2 )
```

```
c(xbar, ybar, sxx, sxy)
```

```
[1] NA NA NA NA
```

Example: Delhi Air Quality data

- Regression coefficients removing missing values

```
x <- sqrt(aqi$no2)
y <- sqrt(aqi$pm10)
xbar <- mean(x, na.rm = TRUE)
ybar <- mean(y, na.rm = TRUE)
sxy <- sum( (x - xbar) * (y - ybar), na.rm = TRUE)
sxx <- sum( (x - xbar)^2, na.rm = TRUE)
```

```
c(xbar, ybar, sxx, sxy)
```

```
[1] 4.703466 11.808792 7506.779092 11173.744082
```

Example: Delhi Air Quality data

- Even this may not be correct: should remove both if only one missing!

```
x <- sqrt(aqi$no2)
y <- sqrt(aqi$pm10)
ok <- is.finite(x) & is.finite(y)
xbar <- mean(x[ok]); ybar <- mean(y[ok])
sxy <- sum( (x[ok] - xbar) * (y[ok] - ybar))
sxx <- sum( (x[ok] - xbar)^2)
```

```
c(xbar, ybar, sxx, sxy)
```

```
[1]      4.700322      11.827206 7480.368670 11173.964771
```

Example: Delhi Air Quality data

- Coefficients of linear regression line

```
bhat <- sxy / sxx  
ahat <- ybar - bhat * xbar  
c(intercept = ahat, slope = bhat)
```

```
intercept    slope  
4.805996    1.493772
```

Example: Delhi Air Quality data

- Coefficients of linear regression line

```
bhat <- sxy / sxx  
ahat <- ybar - bhat * xbar  
c(intercept = ahat, slope = bhat)
```

```
intercept    slope  
4.805996    1.493772
```

- Standard R function to do this

```
lm(sqrt(pm10) ~ sqrt(no2), data = aqi)
```

Call:

```
lm(formula = sqrt(pm10) ~ sqrt(no2), data = aqi)
```


Coefficients:

```
(Intercept)    sqrt(no2)  
      4.806         1.494
```

Example: Delhi Air Quality data

- Add regression line to scatter plot

```
plot(sqrt(pm10) ~ sqrt(no2), data = aqi)  
abline(ahat, bhat, col = "red")
```



Functions

- We have used several functions above
 - `sqrt()`, `mean()`, `sum()` — basic mathematical / summary functions
 - `plot()` — "high level" plotting function
 - `abline()` — "low level" plotting function
 - `lm()` — "high level" modeling function

Functions

- We will discuss graphics functions in more detail later

Functions

- We will discuss graphics functions in more detail later
- `lm()` is a good example to study the behaviour of modeling functions in R

Functions

- We will discuss graphics functions in more detail later
- `lm()` is a good example to study the behaviour of modeling functions in R
- Recall:
 - Functions have one or more arguments
 - All arguments have names (available as variables inside the function)
 - Arguments may be compulsory or optional
 - Optional arguments usually have "default" values

Functions

- `lm()` fits a more general class of models known as *linear models*

```
str(lm)
```

```
function (formula, data, subset, weights, na.action, method = "qr", model = TRUE,  
  x = FALSE, y = FALSE, qr = TRUE, singular.ok = TRUE, contrasts = NULL, offset,  
  ...)
```

Functions

- `lm()` fits a more general class of models known as *linear models*

```
str(lm)
```

```
function (formula, data, subset, weights, na.action, method = "qr", model = TRUE,  
  x = FALSE, y = FALSE, qr = TRUE, singular.ok = TRUE, contrasts = NULL, offset,  
  ...)
```

- These calls are equivalent

```
fm1 <- lm(pm10 ~ no2, aqi, (no2 < 100))  
fm2 <- lm(pm10 ~ no2, data = aqi, method = "qr", subset = (no2 < 100))  
coef(fm1)
```

```
(Intercept)      no2  
52.011396     4.228562
```

```
coef(fm2)
```

```
(Intercept)      no2  
52.011396     4.228562
```

Functions

- Rule: named arguments are matched by name, remaining by position
- Convention:
 - First few "standard" arguments are usually unnamed (matched by position)
 - Usually unnamed arguments are *not* used after named arguments

Functions

- Rule: named arguments are matched by name, remaining by position
- Convention:
 - First few "standard" arguments are usually unnamed (matched by position)
 - Usually unnamed arguments are *not* used after named arguments
- The following call is equivalent to previous two, but not recommended

```
fm3 <- lm(pm10 ~ no2, data = aqi, no2 < 100)  
coef(fm3)
```

```
(Intercept)      no2  
  52.011396    4.228562
```


Functions

- The return value of `lm()` is a list

```
str(fm2)
```

```
List of 13
```

```
$ coefficients : Named num [1:2] 52.01 4.23
..- attr(*, "names")= chr [1:2] "(Intercept)" "no2"
$ residuals    : Named num [1:3810] 128 129 157 143 112 ...
..- attr(*, "names")= chr [1:3810] "1" "2" "3" "4" ...
$ effects      : Named num [1:3810] -9466 3472 155 139 108 ...
..- attr(*, "names")= chr [1:3810] "(Intercept)" "no2" "" "" ...
$ rank         : int 2
$ fitted.values: Named num [1:3810] 132 120 141 103 115 ...
..- attr(*, "names")= chr [1:3810] "1" "2" "3" "4" ...
$ assign       : int [1:2] 0 1
$ qr           :List of 5
..- attr(*, "names")= chr [1:5] "qr" "rank" "fitted" "residuals" "coefficients"
```

Functions

- Individual elements can be extracted using list indexing

```
fm2$coefficients
```

```
(Intercept)      no2  
52.011396      4.228562
```

```
fm2$residuals
```

```
      1      2      3      4      5      6  
127.64591729 129.33160471 157.18879234 143.24585460 111.56016718 131.33160471  
      7      8      9     10     11     12  
 98.01729213 107.93154203 109.16010450 173.93154203  98.61722945 414.61722945  
     13     15     16     17     18     19  
255.07435439 337.41735482 653.33160471  55.58885503 113.64591729  36.27454245  
     20     21     22     23     24     25  
 95.90316761 104.04597998 -2.49689508  73.58885503 -16.72545755 -3.95402002  
     26     28     29     30     31     32  
  9.27454245 17.70297955 19.47441708 27.93154203 26.16010450 11.38866697  
     33     34     35     36     37     38  
13.93154203  9.93154203  3.70297955  4.93154203  1.70297955  4.56016718  
     39     40     41     42     43     44
```

Attributes

- The names (derived from row names of the data) can be used as index

```
fm2$residuals["25"]
```

```
25  
-3.95402
```

- The names associated with a vector can be obtained using `names()`

```
names(fm2$residuals)
```

```
[1] "1"  "2"  "3"  "4"  "5"  "6"  "7"  "8"  "9"  "10" "11" "12"  
[13] "13" "15" "16" "17" "18" "19" "20" "21" "22" "23" "24" "25"  
[25] "26" "28" "29" "30" "31" "32" "33" "34" "35" "36" "37" "38"  
[37] "39" "40" "41" "42" "43" "44" "45" "46" "47" "48" "49" "50"  
[49] "51" "52" "53" "54" "55" "56" "57" "58" "59" "60" "61" "62"  
[61] "63" "64" "65" "66" "67" "68" "69" "70" "71" "72" "73" "74"  
[73] "75" "76" "77" "78" "79" "80" "81" "82" "83" "84" "85" "86"  
[85] "87" "88" "89" "90" "91" "92" "93" "94" "95" "96" "97" "98"  
[97] "99" "100" "101" "102" "103" "104" "105" "106" "107" "108" "109" "110"  
[109] "111" "112" "113" "114" "115" "116" "117" "118" "119" "120" "121" "122"  
[121] "123" "124" "125" "126" "127" "128" "129" "130" "131" "132" "133" "134"
```

Attributes

- These names are actually stored as an *attribute* called "names"

```
attr(fm2$residuals, "names")
```

```
[1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "11" "12"  
[13] "13" "15" "16" "17" "18" "19" "20" "21" "22" "23" "24" "25"  
[25] "26" "28" "29" "30" "31" "32" "33" "34" "35" "36" "37" "38"  
[37] "39" "40" "41" "42" "43" "44" "45" "46" "47" "48" "49" "50"  
[49] "51" "52" "53" "54" "55" "56" "57" "58" "59" "60" "61" "62"  
[61] "63" "64" "65" "66" "67" "68" "69" "70" "71" "72" "73" "74"  
[73] "75" "76" "77" "78" "79" "80" "81" "82" "83" "84" "85" "86"  
[85] "87" "88" "89" "90" "91" "92" "93" "94" "95" "96" "97" "98"  
[97] "99" "100" "101" "102" "103" "104" "105" "106" "107" "108" "109" "110"  
[109] "111" "112" "113" "114" "115" "116" "117" "118" "119" "120" "121" "122"  
[121] "123" "124" "125" "126" "127" "128" "129" "130" "131" "132" "133" "134"  
[133] "135" "136" "137" "138" "139" "140" "141" "142" "143" "144" "145" "146"  
[145] "147" "148" "149" "150" "151" "152" "153" "154" "155" "156" "157" "158"  
[157] "159" "160" "161" "162" "163" "164" "165" "167" "168" "169" "170" "171"  
[169] "172" "173" "174" "175" "176" "177" "178" "179" "180" "181" "182" "183"  
[181] "184" "185" "186" "187" "188" "189" "190" "191" "192" "193" "194" "195"  
[193] "196" "197" "198" "199" "200" "201" "202" "203" "204" "205" "206" "207"  
[205] "208" "209" "210" "211" "212" "213" "214" "215" "216" "217" "218" "219"
```

Attributes

- This is true for all vector objects, including lists

```
attr(fm2, "names")
```

```
[1] "coefficients" "residuals"    "effects"      "rank"         "fitted.values"  
[6] "assign"       "qr"           "df.residual"  "na.action"    "xlevels"  
[11] "call"         "terms"        "model"
```

Attributes

- Attributes are arbitrary R objects that can be attached to any other object
- Typically used for programming convenience, normally not seen by users
- However, some attributes are "special"

Attributes

- The "names" attribute can be extracted using the function `names()`
- `dimnames()` similarly gives row / column names for matrices and arrays

```
dimnames(Titanic)
```

```
$Class  
[1] "1st"  "2nd"  "3rd"  "Crew"  
  
$Sex  
[1] "Male"  "Female"  
  
$Age  
[1] "Child" "Adult"  
  
$Survived  
[1] "No"  "Yes"
```

Attributes

- The "names" attribute can be extracted using the function `names()`
- `dimnames()` similarly gives row / column names for matrices and arrays

```
attr(Titanic, "dimnames")
```

```
$Class  
[1] "1st"  "2nd"  "3rd"  "Crew"  
  
$Sex  
[1] "Male"  "Female"  
  
$Age  
[1] "Child" "Adult"  
  
$Survived  
[1] "No"  "Yes"
```


Attributes

- For example, column names can be obtained as

```
dimnames(Titanic)[[2]]
```

```
[1] "Male"  "Female"
```

- There are convenient shortcuts called `rownames()` and `colnames()`

```
colnames(Titanic)
```

```
[1] "Male"  "Female"
```

Attributes

- In fact, we can easily verify that this is what `colnames()` is doing by printing it

colnames

```
function (x, do.NULL = TRUE, prefix = "col")
{
  if (is.data.frame(x) && do.NULL)
    names(x)
  else dimnames(x)[[2L]] %||% if (do.NULL)
    NULL
  else {
    nc <- NCOL(x)
    if (nc > 0L)
      paste0(prefix, seq_len(nc))
    else character()
  }
}
<bytecode: 0x7fb403a14358>
<environment: namespace:base>
```

- All R functions can be easily inspected in this way

Attributes

- Another very important attribute is "class"
- For example, the return value of `lm()` has class "lm"

```
attr("class")
```

```
[1] "lm"
```

- The class of an object can also be obtained using the function `class()`

```
class(fm2)
```

```
[1] "lm"
```

Attributes

- The class of an object can (usually) be "removed" by setting it to `NULL`
- This is not something you should actually do!

```
class(fm2) <- NULL
```

- Such objects will no longer have a "class" attribute

```
attr(fm2, "class")
```

```
NULL
```

- However, it will still have a class (implicitly)

```
class(fm2)
```

```
[1] "list"
```

Class, generic functions, and methods

- The class of an object is fundamental to how R works
- Every object in R must have a class
- This is true even if the object does not have a class attribute

```
class(colnames)
```

```
[1] "function"
```

```
attr(colnames, "class")
```

```
NULL
```

```
class(Titanic)
```

```
[1] "table"
```

```
attr(Titanic, "class")
```

```
[1] "table"
```

Class, generic functions, and methods

- The main use of the class of an object is in how *generic functions* behave
- Generic functions are intended to perform general tasks, like
 - `print()`
 - `plot()`
 - `summary()`
- But details of what these functions should do depends on the input

Class, generic functions, and methods

```
print(Titanic[, , 1, 1])
```

	Sex	
Class	Male	Female
1st	0	0
2nd	0	0
3rd	35	17
Crew	0	0

```
fm1 <- lm(pm10 ~ no2, aqi, subset = (no2 < 100))  
print(fm1)
```

Call:

```
lm(formula = pm10 ~ no2, data = aqi, subset = (no2 < 100))
```

Coefficients:

(Intercept)	no2
52.011	4.229

Class, generic functions, and methods

```
summary(Titanic)
```

```
Number of cases in table: 2201
```

```
Number of factors: 4
```

```
Test for independence of all factors:
```

```
  Chisq = 1637.4, df = 25, p-value = 0
```

```
Chi-squared approximation may be incorrect
```


Class, generic functions, and methods

```
summary(fm1)
```

Call:

```
lm(formula = pm10 ~ no2, data = aqi, subset = (no2 < 100))
```

Residuals:

Min	1Q	Median	3Q	Max
-235.21	-51.35	-17.91	28.72	830.25

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	52.0114	2.7786	18.72	<2e-16 ***
no2	4.2286	0.1014	41.71	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 83.24 on 3808 degrees of freedom

(118 observations deleted due to missingness)

Multiple R-squared: 0.3136, Adjusted R-squared: 0.3135

F-statistic: 1740 on 1 and 3808 DF, p-value: < 2.2e-16

Class, generic functions, and methods

- But suppose we make a copy of `fm1` and remove the class attribute from it

```
fm2 <- fm1  
class(fm2) <- NULL  
class(fm1)
```

```
[1] "lm"
```

```
class(fm2)
```

```
[1] "list"
```

- `fm1` and `fm2` represent the same model fit
- But the different class means that `print()` and `summary()` behave differently

Class, generic functions, and methods

```
summary(fm2)
```

	Length	Class	Mode
coefficients	2	-none-	numeric
residuals	3810	-none-	numeric
effects	3810	-none-	numeric
rank	1	-none-	numeric
fitted.values	3810	-none-	numeric
assign	2	-none-	numeric
qr	5	qr	list
df.residual	1	-none-	numeric
na.action	118	omit	numeric
xlevels	0	-none-	list
call	4	-none-	call
terms	3	terms	call
model	2	data.frame	list

Class, generic functions, and methods

```
print(fm2)
```

```
$coefficients
```

```
(Intercept)      no2  
  52.011396    4.228562
```

```
$residuals
```

1	2	3	4	5	6
127.64591729	129.33160471	157.18879234	143.24585460	111.56016718	131.33160471
7	8	9	10	11	12
98.01729213	107.93154203	109.16010450	173.93154203	98.61722945	414.61722945
13	15	16	17	18	19
255.07435439	337.41735482	653.33160471	55.58885503	113.64591729	36.27454245
20	21	22	23	24	25
95.90316761	104.04597998	-2.49689508	73.58885503	-16.72545755	-3.95402002

- This kind of customized output is achieved by *methods*

Class, generic functions, and methods

- Methods are specific implementations of a generic function customized to its input
- The appropriate method is chosen by looking at the *class* of the input argument

Class, generic functions, and methods

- Methods are specific implementations of a generic function customized to its input
- The appropriate method is chosen by looking at the *class* of the input argument
- The methods available for a generic function can be obtained using the `methods()` function

```
methods("summary")
```

```
[1] summary.aov                summary.aovlist*
[3] summary.aspell*            summary.check_packages_in_dir*
[5] summary.connection         summary.data.frame
[7] summary.Date               summary.default
[9] summary.ecdf*              summary.factor
[11] summary.glm                summary.hcl_palettes*
[13] summary.infl*              summary.lm
[15] summary.loess*             summary.manova
[17] summary.matrix             summary.mlm*
[19] summary.nls*               summary.packageStatus*
[21] summary.POSIXct            summary.POSIXlt
[23] summary.ppr*               summary.prcomp*
[25] summary.prcomp*            summary.prcomp.time
```

Class, generic functions, and methods

```
methods("print") # similar but much longer list
```

```
[1] print.acf*  
[2] print.activeConcordance*  
[3] print.AES*  
[4] print.anova*  
[5] print.aov*  
[6] print.aovlist*  
[7] print.ar*  
[8] print.Arima*  
[9] print.arima0*  
[10] print.AsIs  
[11] print.aspell*  
[12] print.aspell_inspect_context*  
[13] print.bibentry*  
[14] print.Bibtex*  
[15] print.browseVignettes*  
[16] print.by
```

Class, generic functions, and methods

- All available methods for a given class can be similarly obtained

```
methods(class = "lm")
```

```
[1] add1          alias          anova          case.names     coerce
[6] confint       cooks.distance deviance       dfbeta         dfbetas
[11] drop1         dummy.coef     effects       extractAIC     family
[16] formula       hatvalues     influence     initialize     kappa
[21] labels        logLik        model.frame   model.matrix   nobs
[26] plot          predict        print         proj           qr
[31] residuals     rstandard     rstudent      show           simulate
[36] slotsFromS3   summary        variable.names vcov
see '?methods' for accessing help and source code
```


Class, generic functions, and methods

- The name of a specific method appears to have the form `generic.class`
- However, one should always call the generic function, not the method directly
- This is not OK:

```
summary.lm(fm1)
```

- Instead, use

```
summary(fm1)
```

- In fact, many methods cannot be called directly because they are "hidden"

Class, generic functions, and methods

- This is a form of *Object Oriented Programming* (OOP) in R
- Python also has OOP, but
 - Methods are usually tied to a class, not a *generic* function
 - One notable exception is the `__str__()` method, which is used by `print()`

Getting help

- R has an extensive collection of functions (even more if we include add-on packages)
- It is impossible for anyone to know them all, or remember details
- Fortunately, R also has an excellent help system

Getting help

- R has an extensive collection of functions (even more if we include add-on packages)
- It is impossible for anyone to know them all, or remember details
- Fortunately, R also has an excellent help system
- Every function and dataset in R (and add-on packages) must be documented
- The documentation can be accessed by `help(name)` or `?name`
- For example: `help(seq)`, `help(summary)`, etc.
- A more general (but limited) search can be performed using `help.search("search-string")`

Getting help

- How the help is shown depends on the *interface* being used
- RStudio has a separate help tab (which also allows searching)

Getting help

- How the help is shown depends on the *interface* being used
- RStudio has a separate help tab (which also allows searching)
- However, before using the help system, you should know how methods are documented

Help for generic functions and methods

- Generic functions and methods are distinct functions
- They often have different help pages
- In fact, many add-on packages define new methods for generics in another package
- These are always documented in a separate help page

Help for generic functions and methods

- To get help for the generic function `summary()`, type `help(summary)`
- To get help for the `summary()` method for "matrix" objects, type `help(summary.matrix)`
- To get help for the `summary()` method for "lm" objects, type `help(summary.lm)`
- The first two happen to be the same help page, but the third is different

Help for generic functions and methods

- To get help for the generic function `summary()`, type `help(summary)`
- To get help for the `summary()` method for "matrix" objects, type `help(summary.matrix)`
- To get help for the `summary()` method for "lm" objects, type `help(summary.lm)`
- The first two happen to be the same help page, but the third is different
- This is slightly confusing because you are **not** supposed to call `summary.lm()` directly
- More importantly, there may not actually be a `summary()` method for all classes
- For example, "list" objects are handled by a *fallback* method `summary.default()`

Help for generic functions and methods

- To get help for the generic function `summary()`, type `help(summary)`
- To get help for the `summary()` method for "matrix" objects, type `help(summary.matrix)`
- To get help for the `summary()` method for "lm" objects, type `help(summary.lm)`
- The first two happen to be the same help page, but the third is different
- This is slightly confusing because you are **not** supposed to call `summary.lm()` directly
- More importantly, there may not actually be a `summary()` method for all classes
- For example, "list" objects are handled by a *fallback* method `summary.default()`
- The list of available methods are obtained by `methods("summary")` as shown earlier
- All these should have a corresponding help page

Help for generic functions and methods

- The system we described is called "S3" (short for "S version 3")
- The documentation refers to specific methods implemented using this system as "S3 methods"
- To make things more complicated, there are other systems of defining classes and methods
- We will skip the details of these for now

Replacement Functions

- R functions are not allowed to modify its arguments

Replacement Functions

- R functions are not allowed to modify its arguments
- Consider a function that sets negative inputs to 0

$$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{otherwise.} \end{cases}$$

Replacement Functions

- R functions are not allowed to modify its arguments
- Consider a function that sets negative inputs to 0

$$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{otherwise.} \end{cases}$$

- We will refer to this as the *ReLU* function

The Scalar ReLU Function in R

```
sReLU <- function(u) {  
  if (u < 0) u = 0  
  u  
}
```

The Scalar ReLU Function in R

```
sReLU <- function(u) {  
  if (u < 0) u = 0  
  u  
}
```

```
x <- -5  
sReLU(x)
```

```
[1] 0
```

```
x
```

```
[1] -5
```


The Scalar ReLU Function in R

```
sReLU <- function(u) {  
  if (u < 0) u = 0  
  u  
}
```

```
x <- -5  
y <- sReLU(x)
```

The Scalar ReLU Function in R

```
sReLU <- function(u) {  
  if (u < 0) u = 0  
  u  
}
```

```
x <- -5  
y <- sReLU(x)
```

```
x
```

```
[1] -5
```

```
y
```

```
[1] 0
```

The Scalar ReLU Function in Python

```
def sReLU(u):  
    if u < 0:  
        u = 0  
    return u
```

The Scalar ReLU Function in Python

```
def sReLU(u):  
    if u < 0:  
        u = 0  
    return u
```

```
x = -5  
y = sReLU(x)
```

The Scalar ReLU Function in Python

```
def sReLU(u):  
    if u < 0:  
        u = 0  
    return u
```

```
x = -5  
y = sReLU(x)
```

x

-5

y

0

The Vectorized ReLU Function in Python

```
def vReLU(u):  
    for i in range(len(u)):  
        if u[i] < 0:  
            u[i] = 0  
    return u
```

The Vectorized ReLU Function in Python

```
def vReLU(u):  
    for i in range(len(u)):  
        if u[i] < 0:  
            u[i] = 0  
    return u
```

```
from numpy import *  
x = random.normal(size = 10)  
y = x  
x
```

```
array([-1.2993973 , -1.11465781, -0.31597779, -1.29471625,  0.38419193,  
       0.31589083,  1.17404954,  0.06097466,  0.33193044,  0.62444228])
```

y

```
array([-1.2993973 , -1.11465781, -0.31597779, -1.29471625,  0.38419193,  
       0.31589083,  1.17404954,  0.06097466,  0.33193044,  0.62444228])
```

The Vectorized ReLU Function in Python

```
z = vReLU(y)
```


The Vectorized ReLU Function in Python

```
z = vReLU(y)
```

```
z
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

The Vectorized ReLU Function in Python

```
z = vReLU(y)
```

```
z
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

```
y
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

The Vectorized ReLU Function in Python

```
z = vReLU(y)
```

```
z
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

```
y
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

```
x
```

```
array([0.          , 0.          , 0.          , 0.          , 0.38419193,  
       0.31589083, 1.17404954, 0.06097466, 0.33193044, 0.62444228])
```

The Vectorized ReLU Function in R

```
vReLU <- function(u) {  
  for (i in seq_len(length(u))) {  
    if (u[i] < 0)  
      u[i] <- 0  
  }  
  return(u)  
}
```

The Vectorized ReLU Function in R

```
vReLU <- function(u) {  
  for (i in seq_len(length(u))) {  
    if (u[i] < 0)  
      u[i] <- 0  
  }  
  return(u)  
}
```

```
x <- rnorm(10)  
y <- x
```

The Vectorized ReLU Function in R

```
vReLU <- function(u) {  
  for (i in seq_len(length(u))) {  
    if (u[i] < 0)  
      u[i] <- 0  
  }  
  return(u)  
}
```

```
x <- rnorm(10)  
y <- x
```

x

```
[1] -0.5915339 -0.1111574  1.5757390 -1.5187443 -0.3405556  1.4324608  1.3858898  
[8]  0.8068037 -0.6983363 -0.8147241
```

y

```
[1] -0.5915339 -0.1111574  1.5757390 -1.5187443 -0.3405556  1.4324608  1.3858898  
[8]  0.8068037 -0.6983363 -0.8147241
```

The Vectorized ReLU Function in R

```
z <- vReLU(y)
```

```
z
```

```
[1] 0.0000000 0.0000000 1.5757390 0.0000000 0.0000000 1.4324608 1.3858898 0.8068037
```

```
[9] 0.0000000 0.0000000
```

The Vectorized ReLU Function in R

```
z <- vReLU(y)
```

```
z
```

```
[1] 0.0000000 0.0000000 1.5757390 0.0000000 0.0000000 1.4324608 1.3858898 0.8068037  
[9] 0.0000000 0.0000000
```

```
y
```

```
[1] -0.5915339 -0.1111574 1.5757390 -1.5187443 -0.3405556 1.4324608 1.3858898  
[8] 0.8068037 -0.6983363 -0.8147241
```

```
x
```

```
[1] -0.5915339 -0.1111574 1.5757390 -1.5187443 -0.3405556 1.4324608 1.3858898  
[8] 0.8068037 -0.6983363 -0.8147241
```


Alternative: Direct Assignment

- We can instead use direct assignment to a 'subset'

```
y[y < 0] <- 0
```

Alternative: Direct Assignment

- We can instead use direct assignment to a 'subset'

```
y[y < 0] <- 0
```

```
x
```

```
[1] -0.5915339 -0.1111574  1.5757390 -1.5187443 -0.3405556  1.4324608  1.3858898  
[8]  0.8068037 -0.6983363 -0.8147241
```

```
y
```

```
[1] 0.0000000 0.0000000 1.5757390 0.0000000 0.0000000 1.4324608 1.3858898 0.8068037  
[9] 0.0000000 0.0000000
```

Functional programming

- R generally follows a *functional programming* paradigm
- Among other things, this says that functions should not modify its arguments

Functional programming

- R generally follows a *functional programming* paradigm
- Among other things, this says that functions should not modify its arguments
- This is a key difference between Python and R

Functional programming

- R generally follows a *functional programming* paradigm
- Among other things, this says that functions should not modify its arguments
- This is a key difference between Python and R
- One consequence: R needs an *unusual* approach when it needs to modify objects

Example: Modifying Names of a Data Frame

```
d <- data.frame(1, rnorm(5), rexp(5))  
names(d)
```

```
[1] "X1"      "rnorm.5." "rexp.5."
```

```
d
```

```
  X1  rnorm.5.  rexp.5.  
1  1  0.5137902 2.3629231  
2  1 -1.2638761 0.3281384  
3  1  2.4951138 2.8728996  
4  1 -1.7135097 0.7648846  
5  1  0.2733856 0.5005642
```

- Default names are not very nice

Example: Modifying Names of a Data Frame

- Want to change the names to "Constant", "Normal", "Exponential"
- Possible solution using the `setNames()` function

```
setNames(d, c("Constant", "Normal", "Exponential"))
```

	Constant	Normal	Exponential
1	1	0.5137902	2.3629231
2	1	-1.2638761	0.3281384
3	1	2.4951138	2.8728996
4	1	-1.7135097	0.7648846
5	1	0.2733856	0.5005642

Example: Modifying Names of a Data Frame

- But names of `d` are not modified by this

`d`

```
  X1   rnorm.5.   rexp.5.
1  1  0.5137902  2.3629231
2  1 -1.2638761  0.3281384
3  1  2.4951138  2.8728996
4  1 -1.7135097  0.7648846
5  1  0.2733856  0.5005642
```


Example: Modifying Names of a Data Frame

- But names of `d` are not modified by this

```
d
```

```
  X1  rnorm.5.  rexp.5.  
1  1  0.5137902 2.3629231  
2  1 -1.2638761 0.3281384  
3  1  2.4951138 2.8728996  
4  1 -1.7135097 0.7648846  
5  1  0.2733856 0.5005642
```

- Best we can hope for

```
d <- setNames(d, c("Constant", "Normal", "Exponential"))
```

Example: Modifying Names of a Data Frame

- In fact, the **dplyr** package has a more convenient version of this approach

```
d <- dplyr::rename(d, Constant = X1, Normal = rnorm.5., Exponential = rexp.5.)  
d
```

	Constant	Normal	Exponential
1	1	0.5137902	2.3629231
2	1	-1.2638761	0.3281384
3	1	2.4951138	2.8728996
4	1	-1.7135097	0.7648846
5	1	0.2733856	0.5005642

Example: Modifying Names of a Data Frame

- The classic R Alternative: Replacement Functions

```
names(d) <- c("Constant", "Normal", "Exponential")  
d
```

	Constant	Normal	Exponential
1	1	0.5137902	2.3629231
2	1	-1.2638761	0.3281384
3	1	2.4951138	2.8728996
4	1	-1.7135097	0.7648846
5	1	0.2733856	0.5005642

Replacement Functions

- Other similar examples:

```
y[ y < 0 ] <- 0  
class(x) <- NULL  
d$Normal[ d$Normal < 0 ] <- 0  
attr(x, "name") <- value
```

- Common feature: "complex" expression on the LHS of the assignment